

Jupyter & MSTICPy in security operations and threat hunting



The MSTICPy Team



Ian Hellen

Principal Software
Engineer



Pete Bryan

Senior Security
Researcher



Ashwin Patil

Senior Security
Researcher



**Our
Community**

Many & varied

History of MSTICPy

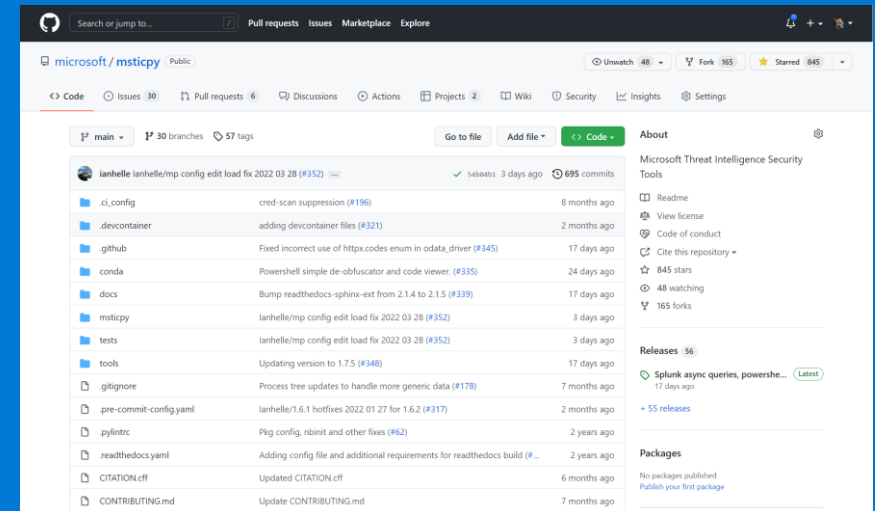
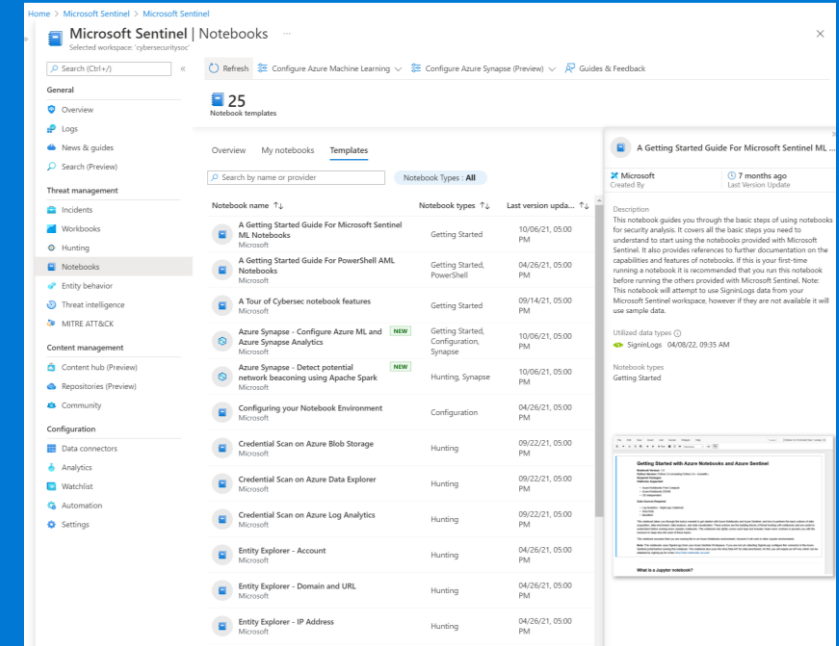
Notebooks in Sentinel

Need to share code with customers

Need to maintain it

Plenty of input and growth

- 150k+ downloads



What's Included

Data Acquisition:

- Sentinel, Kusto, MDE, Graph
- Splunk
- More!

Data Enrichment

- Threat Intel lookups
- Context from Azure APIs
- WhoIs, GeoIP +

Analyzing Data


- Decode
- Extract
- ML

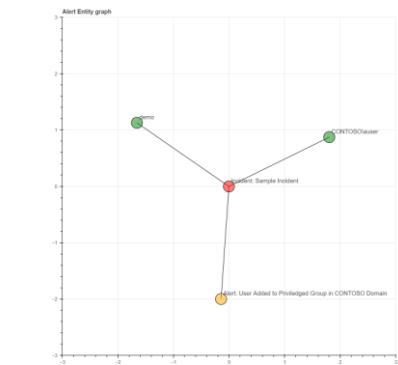
Displaying Data

- Timelines
- Process Trees
- Graphs

```
alerts = qry_prov.SecurityAlert.list_alerts(  
    start='2019-07-21 23:43:18.274492',  
    end='2019-07-27 23:43:18.274492'  
)  
alerts.head()
```

TimeGenerated	AlertDisplayName	Severity	Description
2019-07-22 06:35:13	Suspicious authentication activity	Medium	Although r
2019-07-22 06:35:13	Suspicious authentication activity	Medium	Although r
2019-07-22 07:02:42	Traffic from unrecommended IP addresses was de...	Low	Azure secu
2019-07-26 06:03:16	Traffic from unrecommended IP addresses was de...	Low	Azure secu
2019-07-23 06:42:01	Traffic from unrecommended IP addresses was de...	Low	Azure secu

	IoCType	QuerySubtype	Result	Details
75.137.9	ipv4	None	True	{'pulse_count': 1, 'names': ['Ur
75.137.9	ipv4	None	True	{'verbose_msg': 'IP address in
38.75.137.9	ipv4	None	True	{'score': 1, 'cats': [], 'categoryC
38.75.137.9	ipv4	None	False	0 rows returned.
	ipv4	None	False	Not found.



TimeGenerated	SubjectUserName	SubjectSource	SubjectTarget	NewProcessName	CommentLine	ProcessId	ProcessName	TargetAgentId
2019-07-21 23:43:18.274492	MSTICAlertsWin1	5-15-20	5-15-20	5-15-20	C:\Windows\System32\cmd.exe	5-15-20	C:\Windows\System32\cmd.exe	5-15-20
2019-07-21 23:43:18.274492	MSTICAlertsWin1	5-15-20	5-15-20	5-15-20	C:\Windows\System32\cmd.exe	5-15-20	C:\Windows\System32\cmd.exe	5-15-20
2019-07-21 23:43:18.274492	MSTICAlertsWin1	5-15-20	5-15-20	5-15-20	C:\Windows\System32\cmd.exe	5-15-20	C:\Windows\System32\cmd.exe	5-15-20

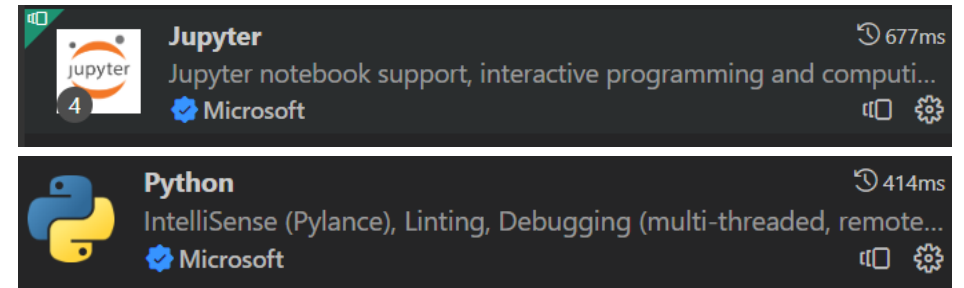
The Agenda



Section	Time
Intros & Setup	20 min
Intro to MSTICPy & Notebooks	30 min
MSTICPy Config	15 min
Break	15 min
Acquiring data with MSTICPy	20 min
Enrichment with MSTICPy	20 min
Jupyter notebooks advanced	15 min
Break	15 min
Data Analysis with MSTICpy	30 min
Data Visualization with MSTICPy	30 min
Putting it into operation	30 min

Setup

- VSCode Installed
 - Jupyter and Python Extensions Installed
- Anaconda Installed
- Azure CLI Installed
- KeyVault Created
- Got API keys for:
 - Alienvault OTX
 - IBM Xforce
 - VirusTotal
 - GeolPLite
- Check you have access –
<https://aka.ms/sentineldemo>



Setup

- Open AnacondaPrompt
- Create new environment
 - ``conda create --name msticpy_training python=3.8``
- Activate Environment
 - ``conda activate msticpy_training``
- Clone the GitHub repo
 - ``git clone https://github.com/microsoft/msticpy-training``
- Installed required packages
 - ``pip install -r msticpy-training\requirements.txt``
- Navigate to our Workshop
 - ``cd msticpy-training\workshops\oct2022``
- Run VSCode from here

Questions & Issues

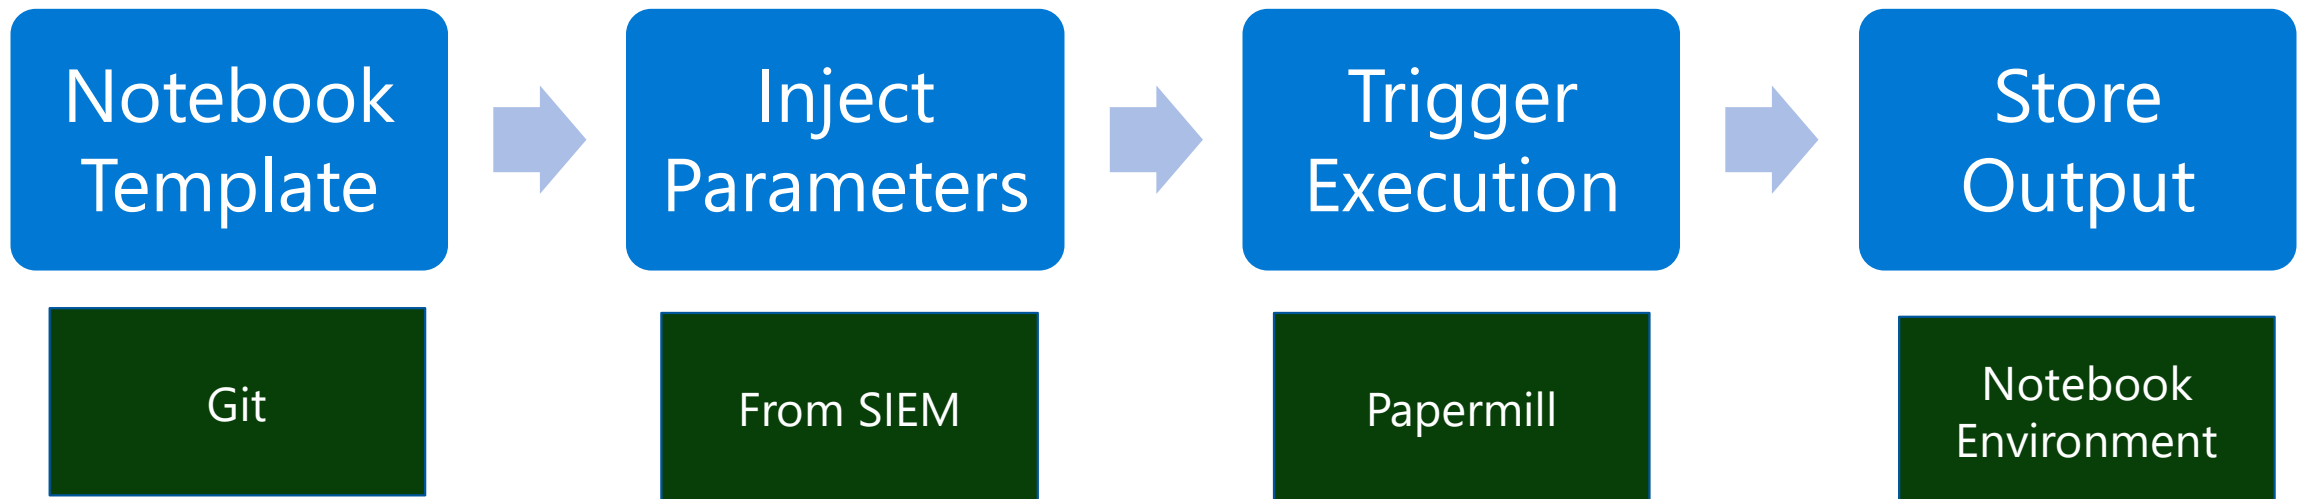
- Teams Channel - https://aka.ms/msticpy_training_teams
- Speak Up
- Breaks to Help Fix Issues

Go to VSCode and select the folder of the msticpy-training repository you just cloned. Open the IntroToMsticpy.ipynb file.

Putting it *All* Together

Operating Model

Automating notebooks execution allows the SOC to benefit from expert knowledge and process



Creating notebook templates

Version Control

Unattended Execution

Execution Options

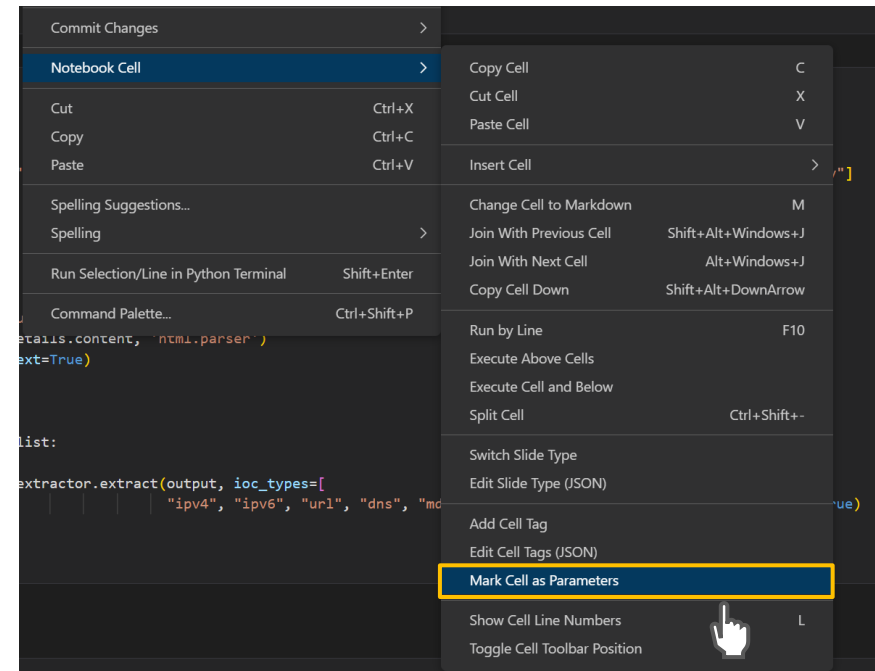
- Papermill parameters
- Default execution path
- Resilient to errors
- Non-interactive

Adding Papermill Parameters

```
1 # papermill default parameters
2 ws_name = "Default"
3 ip_address = ""
4 end = datetime.now(timezone.utc)
5 start = end - timedelta(days=2)
6
```

Create "parameters" cell tag.

Create template cell for parameters.
Some or all values can have defaults.



Allow interactive and automated use

```
1 # papermill default parameters
2 ws_name = "Default"
3 ip_address = ""
4 end = datetime.now(timezone.utc)
5 start = end - timedelta(days=2)
6
```

[]

Notebook parameter cell

```
1 ipaddr_text = nbwidgets.GetText(prompt='Enter the IP Address to search for:', value=ip_address)
2
3 display(ipaddr_text)
4 md("<hr>")
```

[7] ✓ 0.1s

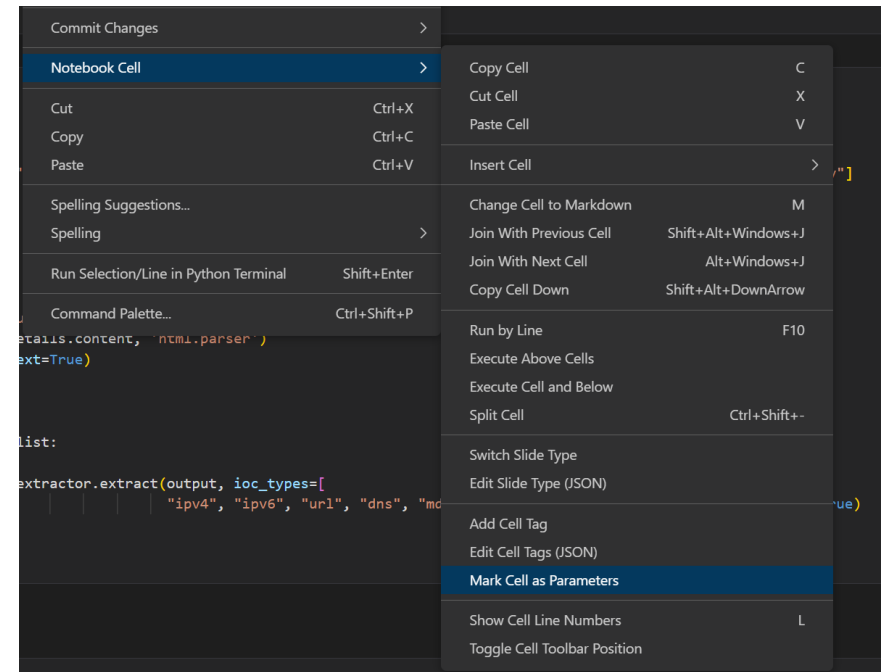
... Enter the IP Address to search for:

Allow editing of parameters in non-blocking UI for interactive use

Exercise – Notebook Parameters

Open the
`AutomatedNotebooks.ipynb`
notebook

Right click on the cell to
parameterize and select
Notebook Cell > Mark Cell as
Parameters



Injecting parameters

1. On the command line

```
$ papermill src/ip_addr.ipynb out/output.ipynb ↵  
-p ip_address "128.1.2.3" ↵  
-p start "2002-07-01 13:05" ↵  
-p end "2002-07-02 13:05" ↵
```

2. In a yaml file

```
ip_address: 128.1.2.3  
start: 2002-07-01 13:05  
end: 2002-07-02 13:05
```

```
$ papermill src/ip_addr.ipynb out/output.ipynb ↵  
-f params.yaml
```

3. From Python

```
return pm.execute_notebook(  
    input_path=input_nb,  
    output_path=output_nb,  
    parameters=nb_params.papermill, # Python dict  
    **nb_kwargs,  
)
```


Exercise – Injecting Parameters

- Open up your Anaconda prompt.
- Attempt to inject parameters into your AutomatedNotebook.ipynb
 - `ip 115.43.212.159`
- Execute the notebook with these parameters and see what output we get

Triggering execution

Scheduled - daily health checks, watch lists

On demand - investigation/analysis tasks

Event triggered - incident/alert triage

You may need **all** of these

Triggering - implementation

Use a cloud service

- Databricks, Azure Synapse, Amazon Sagemaker, etc.
- Likely need to customize for event-triggering

Roll your own

- Cron/Windows job – schedule
- File drop – on demand
- Poller – event-triggered

Build a trigger API

- HTTP endpoint
- JSON parameters

Execution - authentication and secrets

Authentication can be tricky

**Data store (queries)
Services (TI)**

Use a cloud service identity

Store credentials in vault (e.g. Azure Key Vault)

Avoid passing secrets/credentials as Papermill params!!!

Storing and retrieving results

Azure blob

- Cheap!

Output format

- Create output folder structure and naming scheme to organize your outputs

```
/output/2022/08/01/ip-context_124_34_13_59_{UUID}_{date}.ipynb
```

- *Papermill* can strip input code for easier reading
- Create html copies for notebooks with findings (*nbconvert*)

Storing output

Identifying findings: *Interact Scrapbook*

```
1 # Based on results this notebook has a significant finding
2 have_finding = True
✓ 0.7s

1 import scrapbook as sb
2
3 # Surface this as a Scrapbook "scrap"
4 sb.glue("finding", have_finding, display=True)
✓ 0.6s
```

Use scrapbook to check for presence of the scrap

```
1 from pathlib import Path
2 import shutil
3 import scrapbook as sb
4
5 findings_folder = Path("e:/src/blue_team_con/findings")
6 nb_path = Path("e:/src/blue_team_con/scrapbook-test.ipynb")
7 nb = sb.read_notebook(str(nb_path))
8
9 if nb.scrap["finding"].data:
10     if not findings_folder.is_dir():
11         findings_folder.mkdir(parents=True, errors=False)
12         # Copy file (or could create a link)
13         dest_path = shutil.copy(nb_path, findings_folder)
14         print("Copied NB with finding", dest_path)
15
✓ 0.5s

Copied NB with finding e:\src\blue_team_con\findings\scrapbook-test.ipynb
```

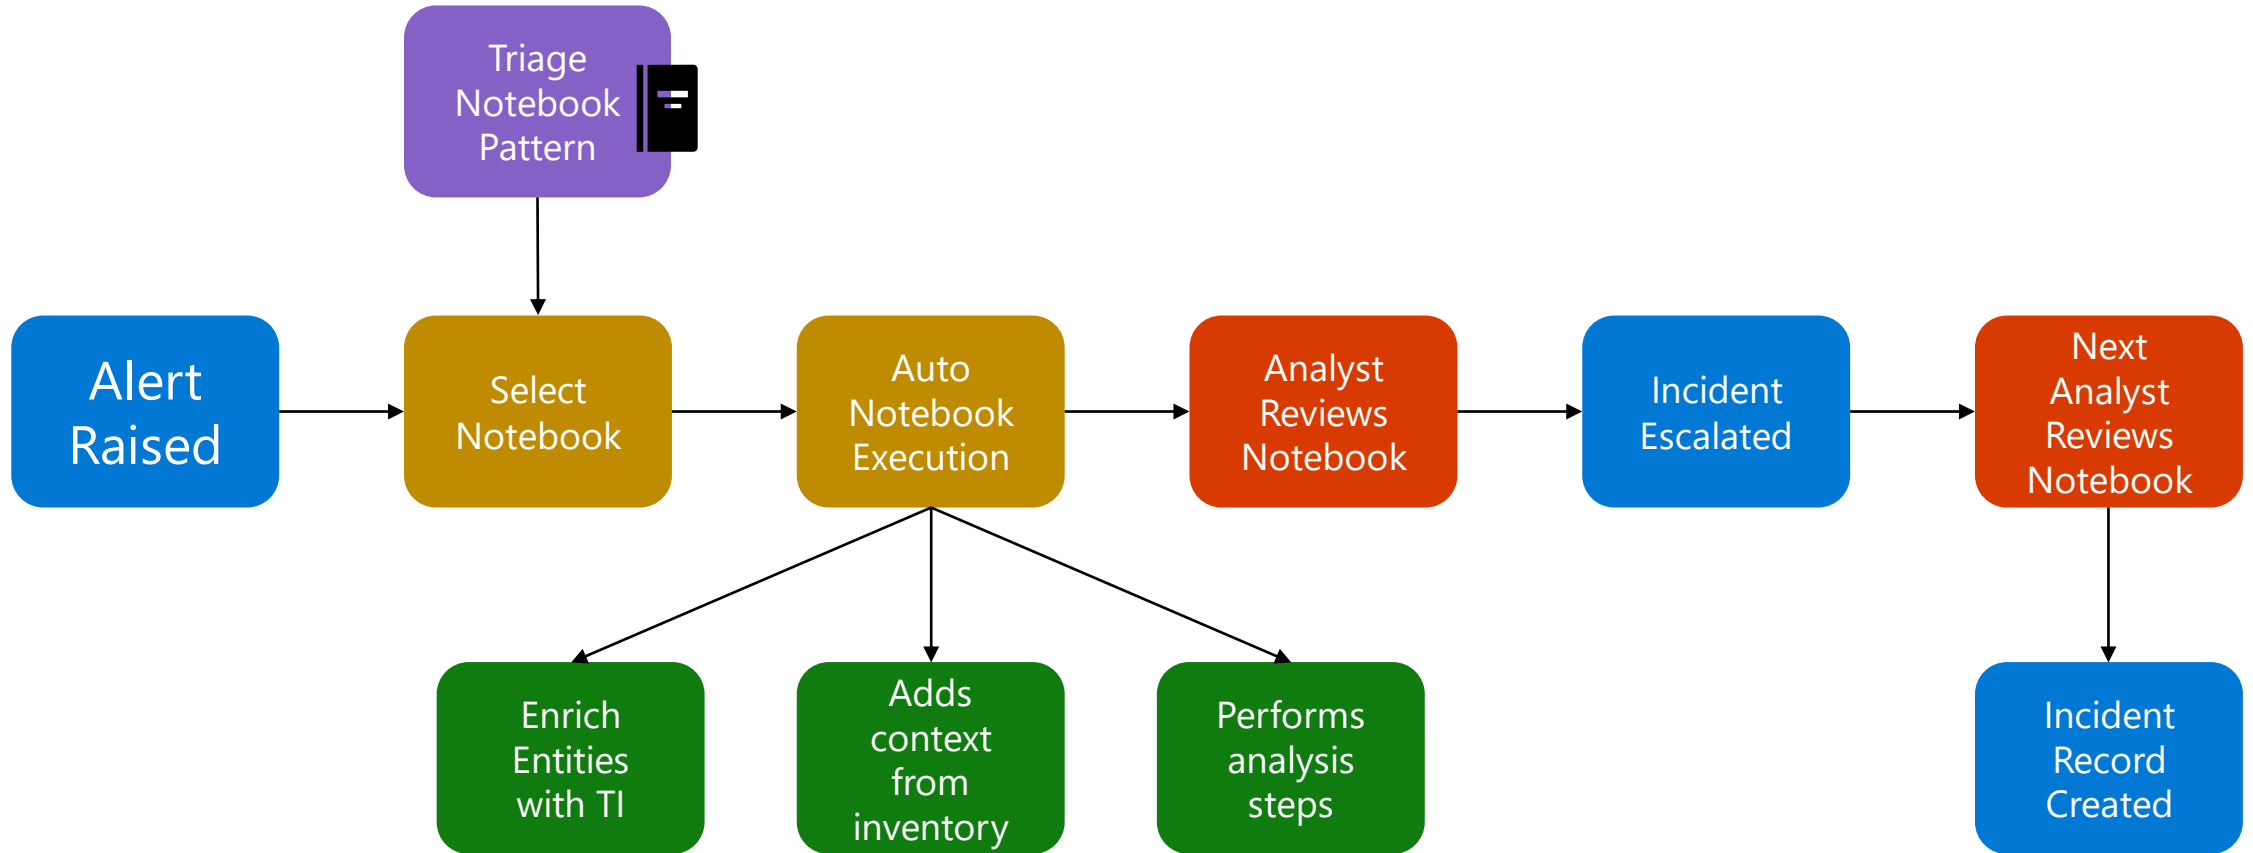
Storing output

Identifying
findings: create
alert/incident

```
1 import msticpy as mp
2 sentinel = mp.MicrosoftSentinel()
3
4 if nb.scraps["finding"].data:
5     sentinel.connect()
6
7     incident_desc = [
8         f"{nb.scraps['finding_desc'].data}",
9         f"Notebook location: {nb_path}"
10    ]
11    sentinel.create_incident(
12        title="Notebook incident created",
13        severity="Medium",
14        status="New",
15        description="\n".join(incident_desc),
16        first_activity_time=datetime.fromtimestamp(nb_path.stat().st_ctime),
17        labels=["notebooks"],
18    )
```

Most incident management systems have
equivalent mechanism

Notebooks for Alert Triage



Notebook automation examples

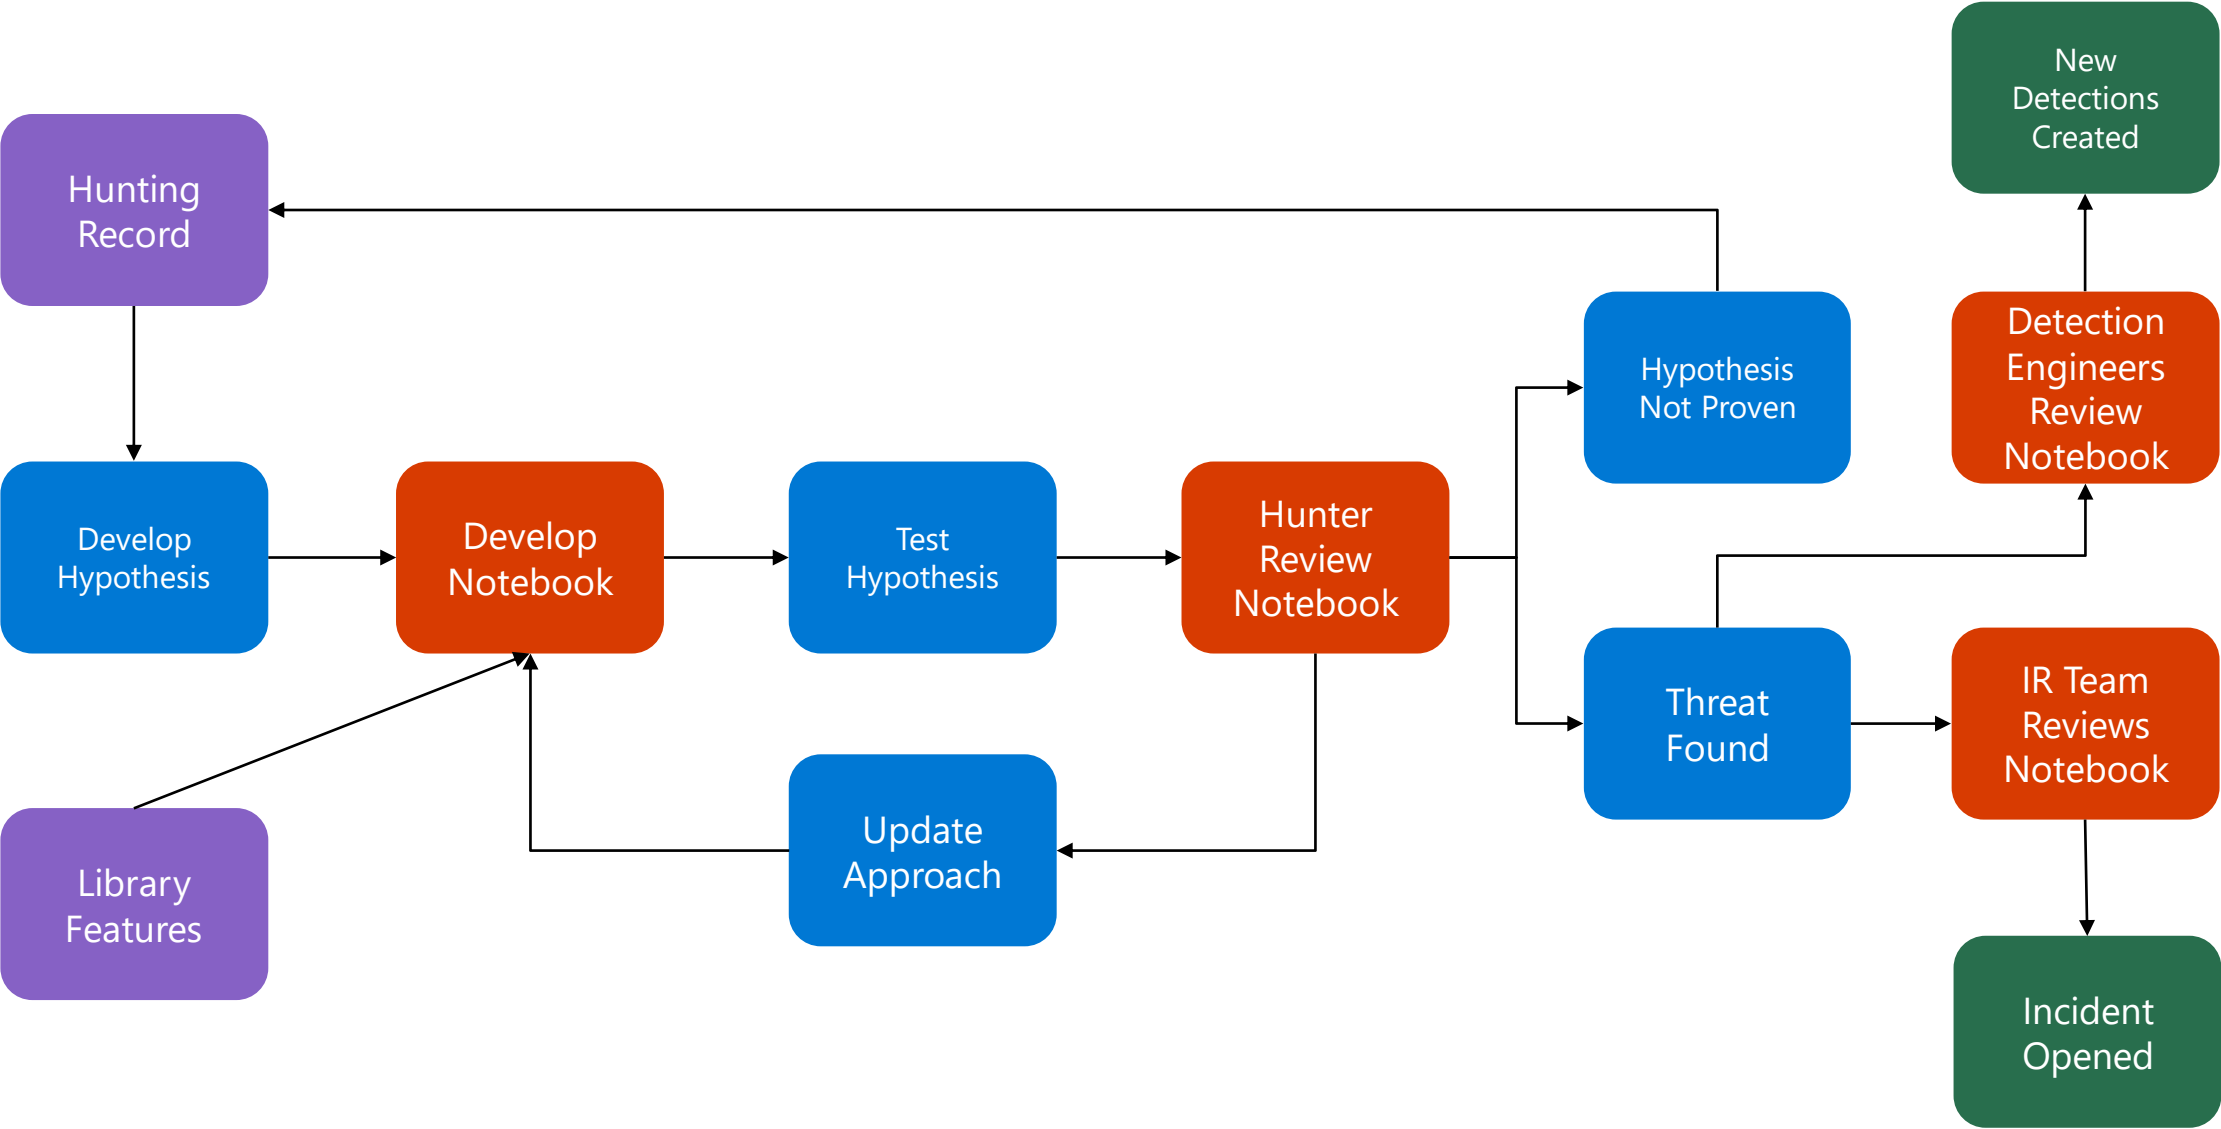
Big brother of the demo

- [Software Defined Monitoring - Using Automated Notebooks and Azure Sentinel to Improve Sec Ops](#)
- Create Azure VM to run notebooks triggered from incidents
- Should be adaptable to other cloud platforms

Our Demo

- Simple solution using Docker + Papermill
- Triggered by YAML parameters file
- Full source on GitHub (see refs)

Notebooks for Threat Hunting



Threat hunting requirements

Usually **ad hoc** but may contain some automated elements

Library support is crucial – make it easy to:

- Query and retrieve information
- Create visualizations
- Repeatable analysis and data extraction/transformation

Package common tasks in parameter-driven notebooks/notebooklets

Apply the same standards as automated notebooks:

- Version control processes (for library and building-block code)
- Output naming and storage

Final Exercise - Optional

- Create your own automated notebook
- Take what you have learnt today and create a notebook using MSTICpy that completes some task
- Parameterize the notebook
- Execute notebook with injected parameters
- Schedule execution for a future time

Find out more

- PyPi

<https://pypi.org/project/msticpy/>

- GitHub Code

<https://github.com/microsoft/msticpy>

- Issues

<https://github.com/microsoft/msticpy/issues>

- Plans

<https://github.com/microsoft/msticpy/discussions>

- ReadTheDocs

<https://msticpy.readthedocs.io/en/latest/index.html>

- msticpy@microsoft.com



Thank You

